

# ITECH2302 Big Data Management

## Laboratory Hadoop pt.2

### Objectives:


- Hadoop file system
- MapReduce

## Activity 1

### Apache Hadoop

#### 1. Start Apache Hadoop

Open a terminal with the ubuntu operating system



Write the following commands:

```
$ ssh localhost  
$ hdfs namenode -format  
$ start-dfs.sh  
$ start-yarn.sh
```

If the output from:

```
$ jps
```

..doesn't look like the following,

5042 DataNode

5299 SecondaryNameNode

4888 NameNode

5516 ResourceManager

5677 NodeManager

6046 Jps

Then maybe the *datanode* didn't start correctly because it was left in a corrupted state. This is easy to fix by using the following commands:

```
$ stop-all.sh
```

```
$ rm -rf /home/goat/hadoopdata/hdfs/datanode/*
```

```
$ start-all.sh
```

Try `jps` again, you should see the datanode listed now.

## Activity 2

### Managing the filesystem

You can create a folder in the Hadoop file system (HDFS) like this:

```
$ hadoop fs -mkdir /data
```

You can copy a file from the ubuntu file system into the Hadoop file system like this:

```
$ hadoop fs -put /home/goat/hadoop_spark/hadoop/lab_data/mapreduce_data/NYSE_DATA.txt /data
```

Now it is available to be operated on by Map/Reduce, or Pig, or Yarn etc. Check that the file was copied over correctly:

```
$ hadoop fs -ls /data
```

You can familiarise yourself with the file system commands if you like, here:

- <https://hadoop.apache.org/docs/r2.4.1/hadoop-project-dist/hadoop-common/FileSystemShell.html>

You can replace `hdfs dfs` with `hadoop fs`

## Activity 3

### Map/Reduce

[https://hadoop.apache.org/docs/r1.2.1/mapred\\_tutorial.html#MapReduce+-+User+Interfaces](https://hadoop.apache.org/docs/r1.2.1/mapred_tutorial.html#MapReduce+-+User+Interfaces)

```
/home/goat/hadoop_spark/hadoop/jar-files/ hadoop-core-1.2.1.jar  
/home/goat/hadoop_spark/hadoop/lab_data/mr/ ProcessUnits.java  
javac -classpath hadoop-core-1.2.1.jar -d units ProcessUnits.java
```

```
javac -classpath /home/goat/hadoop_spark/hadoop/jar-files/hadoop-core-1.2.1.jar -d units  
/home/goat/hadoop_spark/hadoop/lab_data/mr/ProcessUnits.java
```

```
$ jar -cvf units.jar -C units/ .  
$HADOOP_HOME/bin/hadoop fs -mkdir /input_dir  
$HADOOP_HOME/bin/
```

```
hadoop jar units.jar hadoop.ProcessUnits /input_dir output_dir
```

```
hadoop jar units.jar hadoop.ProcessUnits /input_dir output_dir
```

[https://hadoop.apache.org/docs/r1.2.1/mapred\\_tutorial.html](https://hadoop.apache.org/docs/r1.2.1/mapred_tutorial.html)

/home/goat/hadoop\_spark/Hadoop/lab\_data/xxx

/home/goat/hadoop\_spark/Hadoop/lab\_data/mapreduce\_data/NYSE\_mapper.py

[https://www.tutorialspoint.com/map\\_reduce/index.htm](https://www.tutorialspoint.com/map_reduce/index.htm)

**Run examples**

## Activity 4

### Google BigQuery

You might want to check out Google BigQuery:

- <https://cloud.google.com/bigquery>

And its provisions for JSON in its query language:

- [https://cloud.google.com/bigquery/docs/reference/standard-sql/json\\_functions](https://cloud.google.com/bigquery/docs/reference/standard-sql/json_functions)
- [https://docs.snowflake.com/en/sql-reference/functions/parse\\_json.html](https://docs.snowflake.com/en/sql-reference/functions/parse_json.html)